


Note that other plotting methods can be used as desired. For instance, to plot these as a network instead, we can simply extract the adjacency matrix and plot it using standard `igraph` functions

```
net <- igraph::graph.adjacency(MRF_fit$graph, weighted = T, mode = "undirected")
igraph::plot.igraph(net, layout = igraph::layout.circle,
  edge.width = abs(igraph::E(net)$weight),
  edge.color = ifelse(igraph::E(net)$weight < 0,
    'blue',
    'red'))
```



Running CRFs using additional covariates


We can now run a Conditional Random Fields (CRF) model using the provided continuous covariate (`scale.prop.zos`). Again, each species' regression is optimised separately using LASSO regularization. Note that any columns in `data` to the right of column `n_nodes` will be presumed to represent covariates if we don't specify an `n_covariates` argument

```
MRF_mod <- MRFcov(data = Bird.parasites, n_nodes = 4, family = 'binomial')
```

```
## Leave-one-out cv used for the following low-occurrence (rare) nodes:
## Microfilaria ...
## Fitting MRF models in sequence using 1 core ...
```

Visualise the estimated species interaction coefficients as a heatmap. These represent predicted interactions when the covariate is set to its mean (i.e. in this case, when `scale.prop.zos = 0`). If we had included other covariates, then this graph would represent interactions predicted when *all* covariates were set at their means

```
plotMRF_hm(MRF_mod = MRF_mod)
```



Regression coefficients and their relative importances can be accessed as well. This call returns a matrix of the raw coefficient, as well as standardised coefficients (standardised by the `sd` of the covariate). Standardisation in this way helps to compare the relative influences of each parameter on the target species' occurrence probability, but in general the two coefficients will be identical (unless users have not pre-scaled their covariates). The list also contains each variable's relative importance (calculated using the formula $B^2 / \sum(B^2)$, where the vector of B s represents regression coefficients for predictor variables). Variables with an underscore (`_`) indicate an interaction between a covariate and another node, suggesting that conditional dependencies of the two nodes vary across environmental gradients. Because of this, it is recommended to avoid using column names with `_` in them

```
MRF_mod$key_coefs$Hzoosteropis
```

##	Variable	Rel_importance	Standardised_coef	Raw_coef
## 1	Hkillangoi	0.67658708	-2.4508277	-2.4508277
## 5	scale.prop.zos_Microfilaria	0.12292249	-1.0446399	-1.0446399
## 3	Microfilaria	0.09297005	0.9084948	0.9084948
## 4	scale.prop.zos	0.09236272	-0.9055226	-0.9055226
## 2	Plas	0.01196789	-0.3259566	-0.3259566

Finally, a useful capability is to generate some fake data and test predictions. For instance, say we want to know how frequently malaria parasite infections are predicted to occur in sites with high occurrence of microfilaria

```
fake.dat <- Bird.parasites
fake.dat$Microfilaria <- rbinom(nrow(Bird.parasites), 1, 0.8)
fake.preds <- predict_MRF(data = fake.dat, MRF_mod = MRF_mod)
```

The returned object from `predict_MRF` depends on the family of the model. For `family = "binomial"`, we get a list including both linear prediction probabilities and binary predictions (where a linear prediction probability > 0.5 equates to a binary prediction of 1). These binary predictions can be used to estimate parasite prevalence

```
H.zos.pred.prev <- sum(fake.preds$Binary_predictions[, 'Hzoosteropis']) /
nrow(fake.preds$Binary_predictions)
Plas.pred.prev <- sum(fake.preds$Binary_predictions[, 'Plas']) /
nrow(fake.preds$Binary_predictions)
Plas.pred.prev
```

```
## [1] 0.3429844
```

Comparing fits of MRF and CRF models

A key step in the process of model exploration is to determine whether inclusion of covariates actually improves model fit and is warranted when estimating interaction parameters. This is straightforward for binomial models, as we can compare classification accuracy quickly and easily. The `cv_diag` functions will fit models and then determine their predictive performances against the supplied data. We can also compare MRF and CRF models directly to determine whether covariates are warranted. For this dataset, considering that parasite infections are quite rare, we are primarily interested in maximising model Sensitivity (capacity to successfully predict positive infections)

```
mod_fits <- cv_MRF_diag_rep(data = Bird.parasites, n_nodes = 4,
                           n_cores = 1, family = 'binomial', plot = F,
                           compare_null = T,
                           n_folds = 10)
```

```
## Generating node-optimised Conditional Random Fields model
##
## Generating Markov Random Fields model (no covariates)
##
## Calculating model predictions of the supplied data
## Generating CRF predictions ...
## Generating null MRF predictions ...
##
## Calculating predictive performance across test folds
## Processing cross-validation run 1 of 10 ...
## Processing cross-validation run 2 of 10 ...
## Processing cross-validation run 3 of 10 ...
## Processing cross-validation run 4 of 10 ...
## Processing cross-validation run 5 of 10 ...
## Processing cross-validation run 6 of 10 ...
## Processing cross-validation run 7 of 10 ...
## Processing cross-validation run 8 of 10 ...
## Processing cross-validation run 9 of 10 ...
## Processing cross-validation run 10 of 10 ...
```

```
# CRF (with covariates) model sensitivity
quantile(mod_fits$mean_sensitivity[mod_fits$model == 'CRF'], probs = c(0.05, 0.95))
```

```
##          5%          95%
## 0.1245192 0.4227273
```

```
# MRF (no covariates) model sensitivity
quantile(mod_fits$mean_sensitivity[mod_fits$model != 'CRF'], probs = c(0.05, 0.95))
```

```
##          5%          95%
## 0.02714286 0.24358108
```

Bootstrapping the data and running models

We now may want to fit models to bootstrapped subsets of the data to account for parameter uncertainty. Users can change the proportion of observations to include in each bootstrap run with the `sample_prop` option.

```
booted_MRF <- bootstrap_MRF(data = Bird.parasites, n_nodes = 4, family = 'binomial', n_bootstraps
= 10, n_cores = 1)
```

```
## Fitting bootstrap_MRF models in sequence using 1 core...
```

Exploring regression coefficients and interpreting results

Finally, we can explore regression coefficients to get a better understanding of just how important interactions are for predicting species' occurrence probabilities (in comparison to other covariates). This is perhaps the strongest property of CRFs, as comparing the relative importances of interactions and fixed covariates using competing methods (such as Joint Species Distribution Models) is difficult. The `bootstrap_MRF` function conveniently returns a matrix of important coefficients for each node in the graph, as well as their relative importances

```
booted_MRF$mean_key_coefs$Hzosteropis
```

```
##          Variable Rel_importance Mean_coef
## 1          Hkillangoi      0.67748285 -2.4542459
## 5 scale.prop.zos_Microfilaria 0.11969959 -1.0316093
## 3          Microfilaria 0.09421558  0.9152301
## 4          scale.prop.zos 0.08974277 -0.8932410
## 2          Plas      0.01197026 -0.3262276
```

```
booted_MRF$mean_key_coefs$Hkillangoi
```

```
##          Variable Rel_importance Mean_coef
## 1      Hzosteropis      0.78996136 -2.4542459
## 2      Microfilaria      0.12223676 -0.9654195
```

```
## 3 scale.prop.zos      0.08348791 -0.7978605
```

```
booted_MRF$mean_key_coefs$Plas
```

```
##           Variable Rel_importance Mean_coef
## 2      Microfilaria    0.64791745  1.5308224
## 3      scale.prop.zos    0.26641346 -0.9816192
## 4 scale.prop.zos_Microfilaria    0.04831629  0.4180342
## 1           Hzosteropis    0.02942467 -0.3262276
```

```
booted_MRF$mean_key_coefs$Microfilaria
```

```
##           Variable Rel_importance Mean_coef
## 3           Plas    0.35462220  1.5308224
## 4      scale.prop.zos    0.19008774 -1.1207762
## 5 scale.prop.zos_Hzosteropis    0.16104484 -1.0316093
## 2      Hkillangoi    0.14104200 -0.9654195
## 1      Hzosteropis    0.12675845  0.9152301
## 6 scale.prop.zos_Plas    0.02644477  0.4180342
```

Users can also use the `predict_MRFnetworks` function to calculate network statistics for each node in each observation or to generate adjacency matrices for each observation. By default, this function generates a list of `igraph` adjacency matrices, one for each row in data, which can be used to make network plots using a host of other packages. Note, both this function and the `predict_MRF` function rely on data that has a structure exactly matching to the data that was used to fit the model. In other words, the column names and column order need to be identical. The `cutoff` argument is important for binary problems, as this specifies the probability threshold for stating whether or not a species should be considered present at a site (and thus, whether their interactions will be present). Here, we state that a predicted occurrence above 0.33 is sufficient

```
adj_mats <- predict_MRFnetworks(data = Bird.parasites,
                               MRF_mod = booted_MRF,
                               metric = 'eigencentality',
                               cutoff = 0.33)
colnames(adj_mats) <- colnames(Bird.parasites[, 1:4])
apply(adj_mats, 2, summary)
```

```
##           Hzosteropis Hkillangoi   Plas Microfilaria
## Min.           0.0000    0.00000 0.0000    0.00000
## 1st Qu.         0.0000    0.00000 0.0000    0.00000
## Median          0.0000    0.00000 0.0000    0.00000
## Mean            0.1824    0.05791 0.2055    0.08909
## 3rd Qu.         0.4755    0.00000 0.1600    0.00000
## Max.            1.0000    1.00000 1.0000    1.00000
```

Accounting for possible spatial autocorrelation

Lastly, `MRFcov` has the capability to account for possible spatial autocorrelation when estimating interaction parameters. To do this, we incorporate functions from the `mgcv` package to include smoothed Gaussian process spatial regression splines in each node-wise regression. The user must supply a two-column data.frame called `coords`, which will ideally contain Latitude and Longitude values for each row in data. We don't have these coordinates for the `Bird.parasites` dataset, so we will instead create some fake coordinates to showcase the model. Note, these commands were not run here, but feel free to move through them as you did for the above examples

```
Latitude <- sample(seq(120, 140, length.out = 100), nrow(Bird.parasites), TRUE)
Longitude <- sample(seq(-19, -22, length.out = 100), nrow(Bird.parasites), TRUE)
coords <- data.frame(Latitude = Latitude, Longitude = Longitude)
```

The syntax for the `MRFcov_spatial` function is nearly identical to `MRFcov`, with the exception that `coords` must be supplied

```
CRFmod_spatial <- MRFcov_spatial(data = Bird.parasites, n_nodes = 4,
                                  family = 'binomial', coords = coords)
```

Interpretation is also identical to `MRFcov` objects. Here, key coefficients are those that are retained *after* accounting for spatial influences

```
CRFmod_spatial$key_coefs$Hzosteropis
```

Finally, we can compare fits of spatial and non-spatial models just as we did for MRFs and CRFs above

```
cv_MRF_diag_rep_spatial(data = Bird.parasites, n_nodes = 4,  
                          n_cores = 3, family = 'binomial', plot = T, compare_null = T,  
                          coords = coords)
```